

Reference Manual for the Verilog Memory Model Simulation Component

(version 0.4 draft)



Simon Southwell

August 2021
(last updated 20th Sep 2024)

Copyright

Copyright © 2021 – 2024 Simon Southwell (Wyvern Semiconductors)

This document may not, in whole or part, be copied, photocopied, reproduced, translated, or reduced to any electronic medium or machine-readable form without prior written consent from the copyright holder.

Disclaimers

No warranties: the information provided in this document is “as is” without any express or implied warranty of any kind including warranties of accuracy, completeness, merchantability, non-infringement of intellectual property, or fitness for any particular purpose. In no event will the author be liable for any damages whatsoever (whether direct, indirect, special, incidental, or consequential, including, without limitation, damages for loss of profits, business interruption, or loss of information) arising out of the use of or inability to use the information provided in this document, even if the author has been advised of the possibility of such damages.

Simon Southwell (simon@anita-simulators.org.uk)

Cambridge, UK, August 2021



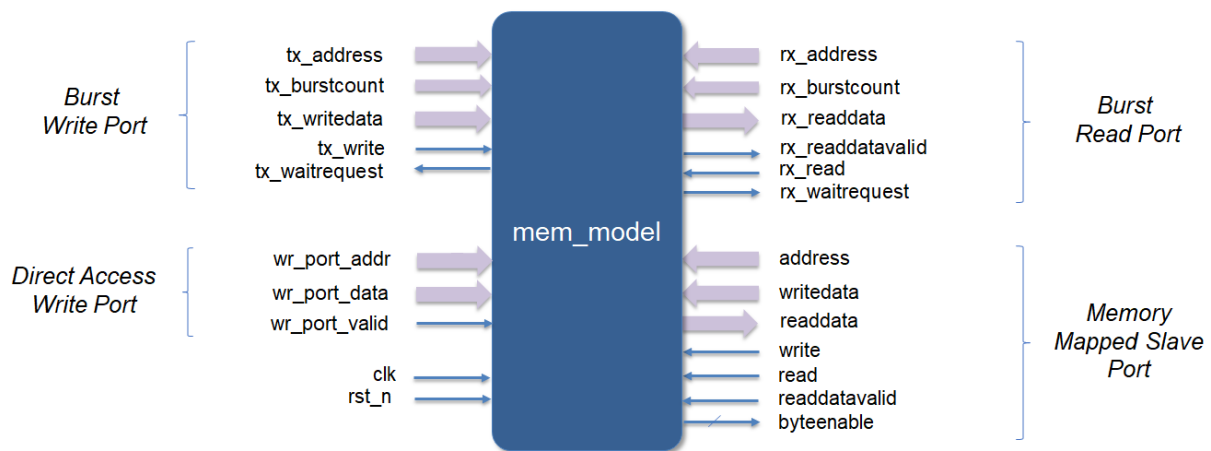
CONTENTS

INTRODUCTION.....	4
INTERNAL MEMORY STRUCTURE	5
COMPILING PLI CODE	6
TEST ENVIRONMENT.....	7
RUNNING THE TEST.....	7
<i>Prerequisites</i>	7
<i>Using make</i>	7
DIRECT ACCESS TO MEMORY FROM C/C++	8
AXI BUS FUNCTIONAL WRAPPER.....	10
REFERENCES.....	11

Introduction

The mem_model component is a Verilog and VHDL simulation test component that allows for a very large memory address space without reserving large amounts of memory, defining large Verilog arrays, or building a truncated memory map into a test bench which could be subject to change in the design. The model uses the logic simulator's programming interfaces to access a C model, pushing the majority of the functionality away from the simulator, make the test bench lightweight, and the memory accesses very fast in simulation compute time.

The component is a lightweight behavioural Verilog module or VHDL component and uses the programming interface to communicate with a set of C/C++ software to implement the actual memory model. The HDL component itself looks like the following:



The component has a clock and reset and four data transfer interfaces.

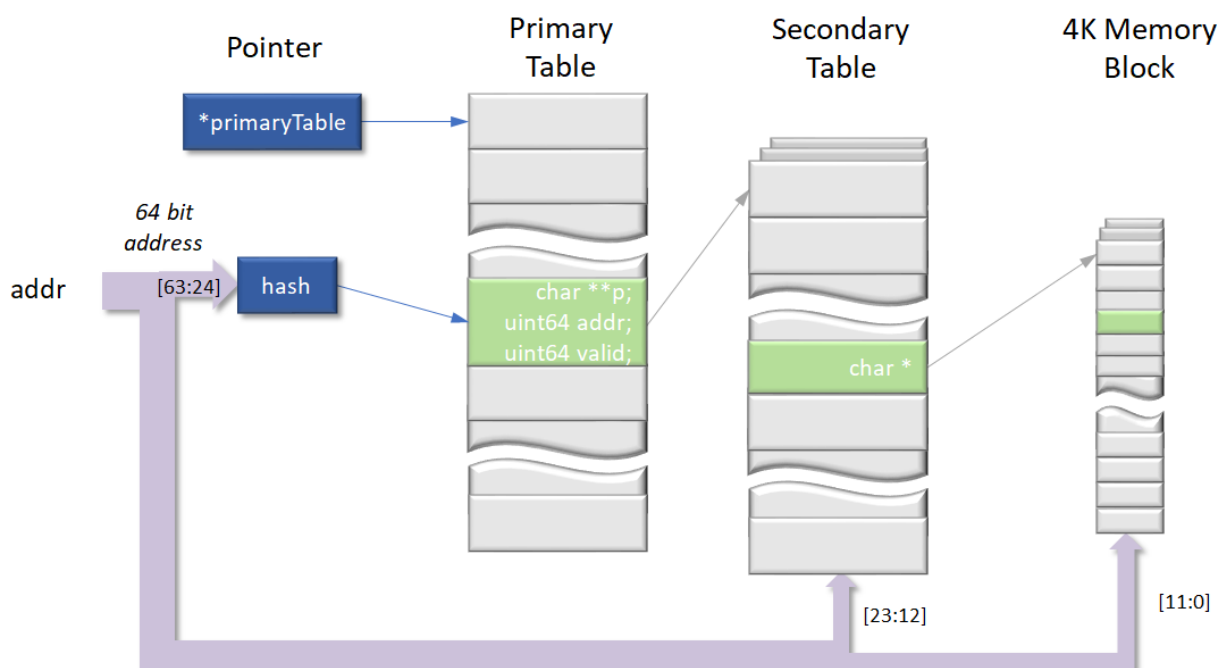
- A memory mapped slave port, for connection to, say, a CPU bus for straight forward word reads and writes with byte enables.
- A burst read port for DMA like read transfers. Note, 4K byte boundaries should not be crossed in a single burst.
- A burst write port for DMA like write transfers. Note, 4K byte boundaries should not be crossed in a single burst.
- A simple SRAM like direct access write port, useful in streaming test data directly to memory.

Not all the interfaces need to be connected, but unused interfaces should have their strobe inputs tied off to 0. If multiple interfaces are active, though, the model will service them in parallel. If two interfaces are accessing the same location in the same cycle, however, behaviour is undefined.

Internal Memory Structure

The memory model component has access to a full 32-bit address space via internal C/C++ memory model software, originally used on a PCIe verification component [1], which actually has capabilities for a full 64-bit address space. It does this with routines defined in `mem.c`, which initialises with no actual memory space allocated. The source file `mem.c` provides two functions for writes and reads—`WriteRamByteBlock()` and `ReadRamByteBlock()`. The user code has access to these functions, as well as some byte and word access hybrid versions (see `src/mem.h`)

The full 64-bit space capability relies on the fact that a simulation cannot possibly write to all 2^{64} locations. Instead, the space is divided into 4K byte chunks which get dynamically allocated as required and are accessed via references in a series of tables which further divided the address space. The starting point for a lookup is the `PrimaryTable`. This table has 4K entries but maps all the top 40 bits of the address space into this space, using a simple hash, XORing the bits in a certain way and then bit reversing the 12-bit result. The `PrimaryTable` entry structure (`PrimaryTable_t`) has a valid field and an address for storing the top 40 bits of the address that hits on the location. If another address upper 40 bits hashes to the same location, then the index pointing to the table entry is simply incremented until an empty entry is found, or we searched the whole table (an error condition). This is summarised in the diagram below.



The primary table entry also contains a pointer to a pointer, which references a secondary table, dynamically allocated when first written to. The secondary tables sub-divide the address space of the lower 24 bits of the address into the 4K byte blocks required. The upper 12 bits of the lower address index into the secondary table, whose entry points to a 4K byte block of memory, dynamically allocated on first access.

Reading from a location simply involves traversing the table. The top 40 bits of the read address are hashed, and index into the primary table. The Primary table entry address is compared with the read address 40 bits and, if different, the index is incremented until a match, an invalid address is encountered, or the whole table is searched. The last two cases are an error condition. The secondary table is then accessed with the next 12 bits and (if pointing to a valid byte block), the lower 12 bits used to index into the page and retrieve the data. At any

point in the traverse, an unallocated table entry of byte block is considered a fatal error—it is not legal to access locations that have not been written.

Compiling PLI code

The model's source code is a set of C files that must be compiled into a shared object, and entries added to a PLI table in order to add the necessary functions in the Verilog domain. The exact procedure varies somewhat between simulators, but examples for Altera's Questa are given here for Verilog. In a file, `veriusertfs.c`, the following code must be present

```
#include "veriusertfs.h"
#include "vpi_user.h"

s_tfcell veriusertfs[] =
{
    {usertask, 0, NULL, 0, MemRead, NULL, "$memread", 1},
    {usertask, 0, NULL, 0, MemWrite, NULL, "$memwrite", 1},
    {0}
};

p_tfcell bootstrap ()
{
    return veriusertfs;
}
```

This table can only be defined once, so if there are other entries required from other PLI code, then these must be combined into a single table. For convenience, the table entries are defined in `mem.h` as a `#define` of `MEM_MODEL_TF_TBL`.

Various flags are required when compiling the code to generate a shared object, load the simulator's PLI library and ensure no (as yet) unreferenced entries are removed at link time. The verification environment (see next section) has example compilation code in a `makefile`, using the virtual processor (VProc) component [2], and calling its `makefile`. The VProc repository has several examples of compiling a shared object, suitable for loading to various simulators, and the reader should reference these for more details.

For Questa, once the shared object is correctly compiled it can be loaded when running `vsim` by using the `-pli <mysharedobj>.so` command line option.

Test Environment

A Verilog test environment for the model is provided. It is dependent on VProc [2], which should be checked out into the same directory as the location of the memory model's repository. Currently it only supports Questa.

The test bench folder (test) contains the following files:

- `tb.v`: the top-level test bench Verilog
- `cpu.v`: a VProc wrapper block that instantiates VProc, and drives the memory mapped master bus
- `files_vlog.tcl`: a list of all the Verilog files required by `vlog`.
- `makefile`: the compilation and execution make file
- `cleanvlib.do`, `compile.do`: ModelSim files for cleaning the Verilog work library and compiling the code, respectively.
- `sim.do`, `simg.do`, `simlog.do`: ModelSim files for running the simulation in various modes (batch, GUI, with logging).

The test source code for running on the Virtual Processor is in a sub-folder `src`. As well as the top level `VUserMain0` code (the virtual processor is set as node 0), it has some driver code for generating the access traffic on the master bus. This driver code is in `mem_vproc_api.cpp` and `mem_vproc_api.h`. The code provides some simple functions to read and write byte, half-word (16 bit) and word (32 bit) values. It is from these functions that calls to the VProc API are made, hiding the details of these, and also managing the byte enable settings.

Running the Test

Prerequisites

The provided `makefile` allows for compilation and execution of the example simulation test. As mentioned above it is only for Questa at this time and has some prerequisites before being able to be run. The following tools are expected to be installed

- Questa Intel Starter Edition. Version used is 2011.1, bundled with Quartus Prime 22.3 Lite Edition. The `MODEL_TECH` environment variable must be set to point to the simulator binaries directory
- `MSYS2/mingw-w64`, including the gcc toolchain and make
- VProc: <https://github.com/wyvernSemi/vproc>, checked out to `vproc/` in the same directory as the `mem_model` folder.

The versions mentioned above are those that have been used to verify the environment, but different versions will likely work just as well. If `MODEL_TECH` environment variable is set correctly, then the `makefile` should just use whichever versions are installed. Note that the test code must be compiled for 64-bit, which the VProc `makefile` does. This is required for Quartus PLI code and may be different for other simulators. So, if compiling PLI code independently of the VProc `makefile` then remember to add the `-m64` flag.

Using make

All the code can be compiled with `make`, and various modes of execution can also be controlled from here. Typing `make help` displays the following usage message:

make help	Display this message
make	Build C/C++ code without running simulation
make compile	Build HDL code without running simulation
make run/sim	Build and run batch simulation
make rungui/gui	Build and run GUI simulation
make runlog/log	Build and run batch simulation with signal logging
make waves	Run wave view in free starter ModelSim (to view runlog signals)
make clean	clean previous build artefacts

Giving no arguments compiles the C/C++ code. The compile argument compiles the Verilog. The various execution modes are for batch, GUI and batch with logging. The batch run with logging assumes there is a wave.do file (generated from the waveform viewer) that it processes into a batch.do with signal logging commands. When run, a vsim.wlf file is produced by the simulator, just as it would when running in GUI mode, with data for the logged signals. The make wave command can then be used to display the logged data without running a new simulation. Note that the batch.do generation is fairly simple at this point, and if the wave.do is constructed with complex mappings the translation may not work. The make clean command deletes all the artefacts from previous builds to ensure the next compilation starts from a fresh state.

When running the tests the simulation should display some access message and print a PASS message at the end. If a comparison fails a ***FAIL*** message is printed. This test code does not exhaustively test the model, with only the slave memory mapped interface being tested. The underlying C/C++ model, however, has been used in the PCIe model [1] extensively, and these test are not meant to repeat that verification. The other interfaces on the component use the same access methods as the slave interface (without the complication of byte enables) and have been used in other verification environments successfully. The test is meant to serve as an example usage within a simulation test environment using, in this case, a VProc [2] based CPU component.

Direct Access to Memory from C/C++

The test environment example discussed above is all that is needed to access memory with bus transactions—in this case from the VProc based component. Another advantage of having the memory model in C, running alongside the test code running in the virtual processor, all of which is compiled into a single package, is that the VProc program can have access to the memory model C code without the need to issue a simulation transaction. This could be advantageous for, say, loading a large set of test data into memory, or for extracting and processing captured data sent to memory from the simulation. This consumes no simulations cycles and runs at the normal speed of an executable program running on the host PC or workstation.

By including the mem.h header in the VProc test code, an API is presented to read and write data to the underlying C model. The prototypes for the API functions are shown below:


```

void      WriteRamByte      (const uint64_t addr, const uint32_t data,
                             const uint32_t node);

void      WriteRamHWord     (const uint64_t addr, const uint32_t data, const int le,
                             const uint32_t node);

void      WriteRamWord      (const uint64_t addr, const uint32_t data, const int le,
                             const uint32_t node);

void      WriteRamDWord     (const uint64_t addr, const uint64_t data, const int le,
                             const uint32_t node);

uint32_t  ReadRamByte       (const uint64_t addr, const uint32_t node);
uint32_t  ReadRamHWord      (const uint64_t addr, const int le, const uint32_t node);
uint32_t  ReadRamWord       (const uint64_t addr, const int le, const uint32_t node);
uint64_t  ReadRamDWord      (const uint64_t addr, const int le, const uint32_t node);

void      WriteRamByteBlock (const uint64_t addr, const PktData_t* const data,
                             const int fbe, const int lbe, const int length,
                             const uint32_t node);

int       ReadRamByteBlock  (const uint64_t addr, PktData_t* const data,
                             const int length, const uint32_t node);

```

The first four functions allow the writing of byte, half-word (16), word (32) and double word (64) values directly to memory. An address and data arguments are given and, for the access greater than a byte, and 'little endian' flag (le) to indicate whether to access little endian (when non-zero) or big-endian (when zero).

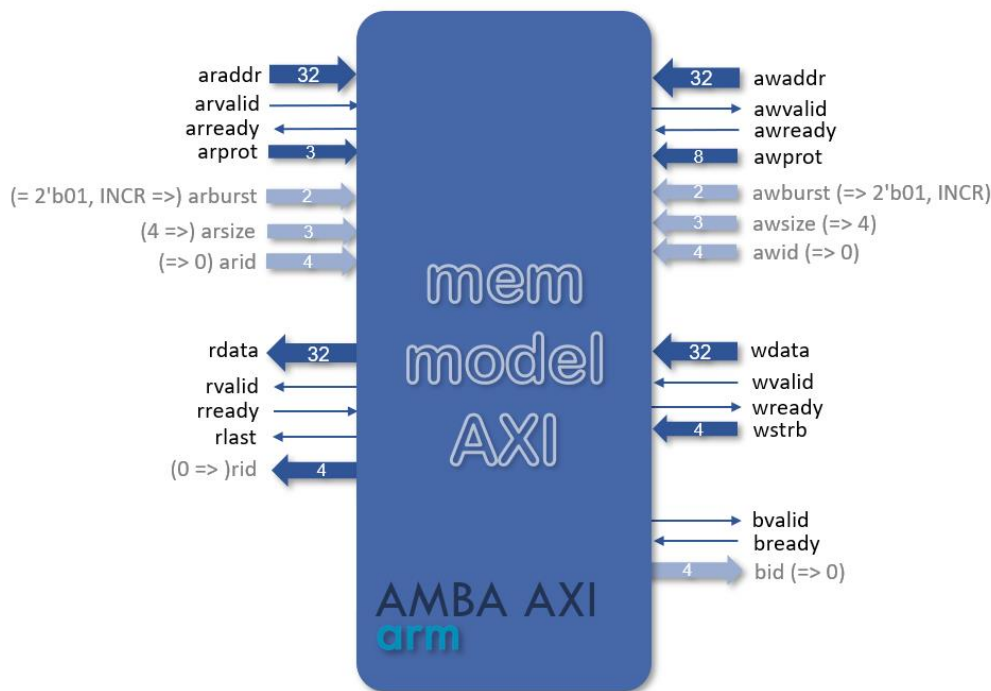
The node argument (for this API) is more like a handle for multiple instantiations of the model and need not necessarily be tied to a VProc node. The expected usage would be to either have a one-to-one connection to a VProc if, say, modelling a multiprocessor system, with each CPU having its own memory. A more likely scenario, though, is that the memory is shared between processors. In this case a test bench would instantiate multiple memory components, one for each processor, but the calling PLI C code uses a common node number. This is, in fact, the default for the mem_model.[ch] code. The Verilog components then act as separate ports into the same memory sub-system. The direct access API, then, becomes a means for shared memory communication between code running on multiple VProc models.

The next four functions are the equivalent for reading from memory, the difference being that the read data is returned by the function. Note that for all the functions the address is a 64-bit data type. As mentioned before, this model, taken from a PCIe root complex model [1], can simulate a full 64-bit address space, even if the Verilog component is limited to 32-bit addresses.

The two block access functions would not normally be used directly (the previously mentioned functions call these), but they can be used if necessary. These take an address, as before, and then a pointer to an array of PktData_t—essentially byte values—where data is taken from, for writes, or data placed into, for reads. The fbe and lbe arguments are byte enables for the first and last 32-bit words, being 4-bit bitmaps. The byte enables are expected to be contiguous for the whole block, so fbe can only be 0xf, 0xe, 0xc and 0x8, whilst the lbe can only be 0x1, 0x3, 0x7 and 0xf, though lbe can be 0 if there is only 4 bytes or less. The byte enables between the first and last words is implied as 0xf. Following the byte enables is a length argument, specifying the number of bytes to be transferred.

AXI Bus Functional Model Wrapper

The default memory model interfaces are Altera Avalon memory mapped bus compatible. Some Verilog and VHDL bus functional model (BFM) wrapper components are also supplied with AXI subordinate interfaces—`mem_model_axi.v` for Verilog and `mem_model.vhd` for VHDL. These support byte enable writes and read and write burst transfers. The diagram below shows the AXI component.



All the ports required for an AXI subordinate component are present. However, at the present time, some are fixed or not used. The `awburst`, `awsize` and `awid` inputs for the write address port only support incrementing bursts, with transfers the full size of the bus with an ID of 0. Similarly, the write response `bid` output is always 0. For the read address port, the `arburst`, `arsize` and `arid` inputs must be as for the write address ports. For the VHDL component, these fixed inputs are the default values and so do not need to be connected and be set externally to these required values. For the read data port, the `rid` output is fixed at 0. Future versions of the model will make use of these inputs.

References

- [1] *PCIe Virtual Host Model Test Component*,
<https://github.com/wyvernSemi/pcievhost/blob/master/doc/pcieVHost.pdf>, Simon Southwell, March 2017.
- [2] *Virtual Processor (VProc)*, <https://github.com/wyvernSemi/vproc> , Simon Southwell, June 2010